

Outline

DSM parameters

- A taxonomy of DSM parameters
- Context type & size
- Feature scaling
- Measuring distance
- Dimensionality reduction
- Examples

Building a DSM

- Sparse matrices
- Example: a verb-object DSM

Distributional Semantic Models

Part 2: The parameters of a DSM

Stephanie Evert¹

with Gabriella Lapesa², Alessandro Lenci³, and Marco Baroni⁴

¹Friedrich-Alexander-Universität Erlangen-Nürnberg, Germany


²Heinrich-Heine-Universität Düsseldorf & GESIS, Germany

³University of Pisa, Italy

⁴ICREA / Universitat Pompeu Fabra Barcelona, Italy

<http://wordspace.collocations.de/doku.php/course:start>

Copyright © 2009–2026 Evert, Lapesa, Lenci & Baroni | Licensed under CC-by-sa version 3.0



General definition of DSMs

A **distributional semantic model** (DSM) is a scaled and/or transformed co-occurrence matrix $\mathbf{M}$, such that each row $\mathbf{x}$ represents the distribution of a target term across contexts.

	get	see	use	hear	eat	kill
knife	0.027	-0.024	0.206	-0.022	-0.044	-0.042
cat	0.031	0.143	-0.243	-0.015	-0.009	0.131
dog	-0.026	0.021	-0.212	0.064	0.013	0.014
boat	-0.022	0.009	-0.044	-0.040	-0.074	-0.042
cup	-0.014	-0.173	-0.249	-0.099	-0.119	-0.042
pig	-0.069	0.094	-0.158	0.000	0.094	0.265
banana	0.047	-0.139	-0.104	-0.022	0.267	-0.042

Term = word, lemma, phrase, morpheme, word pair, ...

General definition of DSMs

Mathematical notation:

- ▶ $k \times n$ co-occurrence matrix $\mathbf{M} \in \mathbb{R}^{k \times n}$ (example: 7×6)
 - ▶ k rows = **target** terms
 - ▶ n columns = **features** or other **dimensions**

$$\mathbf{M} = \begin{bmatrix} m_{11} & m_{12} & \cdots & m_{1n} \\ m_{21} & m_{22} & \cdots & m_{2n} \\ \vdots & \vdots & & \vdots \\ m_{k1} & m_{k2} & \cdots & m_{kn} \end{bmatrix}$$

- ▶ distribution vector $\mathbf{m}_i = i$ -th row of $\mathbf{M}$, e.g. $\mathbf{m}_3 = \mathbf{m}_{\text{dog}} \in \mathbb{R}^n$
- ▶ components $\mathbf{m}_i = (m_{i1}, m_{i2}, \dots, m_{in}) = \text{features of } i\text{-th term:}$

$$\begin{aligned} \mathbf{m}_3 &= (-0.026, 0.021, -0.212, 0.064, 0.013, 0.014) \\ &= (m_{31}, m_{32}, m_{33}, m_{34}, m_{35}, m_{36}) \end{aligned}$$

Term-term matrix

Term-term matrix records co-occurrence frequencies with feature terms for each target term

$\mathbf{m}_{\text{dog}}$ = collocational profile of *dog* ($\approx$ word sketch)

		breed	tail	feed	kill	important	explain	likely
$\mathbf{M} = \begin{bmatrix} \dots & \mathbf{m}_1 & \dots \\ \dots & \mathbf{m}_2 & \dots \\ & \vdots & \\ & \vdots & \\ \dots & \mathbf{m}_k & \dots \end{bmatrix}$	cat	83	17	7	37	—	1	—
	dog	561	13	30	60	1	2	4
	animal	42	10	109	134	13	5	5
	time	19	9	29	117	81	34	109
	reason	1	—	2	14	68	140	47
	cause	—	1	—	4	55	34	55
	effect	—	—	1	6	60	35	17

```
> TT <- DSM_TermTerm
> head(TT, Inf) # extract full co-oc matrix from DSM object
```

Term-context matrix

Term-context matrix records frequency of term in each individual context unit (e.g. document, tweet, encyclopaedia article)

$\mathbf{f}_{\text{dog}}$ = texts related to or mentioning dogs

		Felidae	Pet	Feral	Bloat	Philosophy	Kant	Back pain
$\mathbf{F} = \begin{bmatrix} \dots & \mathbf{f}_1 & \dots \\ \dots & \mathbf{f}_2 & \dots \\ & \vdots & \\ & \vdots & \\ \dots & \mathbf{f}_k & \dots \end{bmatrix}$	cat	10	10	7	—	—	—	—
	dog	—	10	4	11	—	—	—
	animal	2	15	10	2	—	—	—
	time	1	—	—	—	2	1	—
	reason	—	1	—	—	1	4	1
	cause	—	—	—	2	1	2	6
	effect	—	—	—	1	—	1	—

```
> TC <- DSM_TermContext
> head(TC, Inf)
```

Outline

DSM parameters

A taxonomy of DSM parameters

Context type & size

Feature scaling

Measuring distance

Dimensionality reduction


Examples

Building a DSM


Sparse matrices

Example: a verb-object DSM

Building a distributional model



Building a distributional model



Definition of target and feature terms

- ▶ Choice of linguistic unit (targets $\neq$ features)
 - ▶ words
 - ▶ bigrams, trigrams, ...
 - ▶ multiword units, named entities, phrases, ...
 - ▶ morphemes
 - ▶ word pairs (👉 analogy tasks)
- ▶ Mapping to target/feature terms (→ linguistic annotation)
 - ▶ word forms (minimally requires tokenisation)
 - ▶ often lemmatisation or stemming to reduce data sparseness:
go, goes, went, gone, going → *go*
 - ▶ POS disambiguation (*light*/N vs. *light*/A vs. *light*/V)
 - ▶ word sense disambiguation (*bank*_{river} vs. *bank*_{finance})
 - ▶ abstraction: POS tags (or *n*-grams of POS tags) as features

👉 What is the effect of these choices?

Effects of term mapping

Nearest neighbours of *walk* (BNC)

word forms	lemmatised + POS
▶ stroll	▶ hurry
▶ walking	▶ stroll
▶ walked	▶ stride
▶ go	▶ trudge
▶ path	▶ amble
▶ drive	▶ wander
▶ ride	▶ walk (noun)
▶ wander	▶ walking
▶ sprinted	▶ retrace
▶ sauntered	▶ scuttle

<http://clic.cimec.unitn.it/infomap-query/>

Effects of term mapping

Nearest neighbours of *arrivare* (Repubblica)


word forms	lemmatised + POS
▶ giungere	▶ giungere
▶ raggiungere	▶ aspettare
▶ arrivi	▶ attendere
▶ raggiungimento	▶ arrivo (noun)
▶ raggiunto	▶ ricevere
▶ trovare	▶ accontentare
▶ raggiunge	▶ approdare
▶ arrivasse	▶ pervenire
▶ arriverà	▶ venire
▶ concludere	▶ piombare

<http://clic.cimec.unitn.it/infomap-query/>

Selection of target and feature terms

- ▶ Full-vocabulary models are often unmanageable
 - ▶ 762,424 distinct word forms in BNC / 605,910 lemmata
 - ▶ large Web corpora have > 10 million distinct word forms
 - ▶ low-frequency targets (and features) are not reliable (“noisy”)
- ▶ Frequency-based selection
 - ▶ corpus frequency $f \geq F_{\min}$ or n_w most frequent terms
 - ▶ sometimes upper threshold for features: $F_{\min} \leq f \leq F_{\max}$
- ▶ Relevance-based selection of features
 - ▶ criterion from information retrieval: document frequency df (high $df \rightarrow$ uninformative / low $df \rightarrow$ too sparse to be useful)
 - ▶ alternatives: entropy H or chi-squared statistic χ^2
- ▶ Other criteria
 - ▶ POS-based filter: no function words, only verbs, nouns, ...
 - ▶ general dictionary, words required for particular task, ...


Building a distributional model



Term-context matrix: choice of context unit

- ▶ Features are usually **tokens** of the selected context unit, i.e. individual instances of a
 - ▶ document, novel, Wikipedia article, Web page, ...
 - ▶ paragraph, sentence, tweet, ...
 - ▶ “co-occurrence” f_{ij} = frequency of term i in context token j
- ▶ Similar context tokens can be **aggregated**, e.g.
 - ▶ feature = cluster of near-duplicate documents
 - ▶ feature = syntactic structure of sentence (ignoring content)
 - ▶ feature = all tweets from same author (“supertweet”)
 - ▶ f_{ij} = pooled frequency count for aggregate j
- ▶ Generalization: context **types**
 - ▶ e.g. pattern of POS tags around target word
 - ▶ e.g. subcategorisation pattern of target verb

Building a distributional model



Term-term matrix: definition of co-occurrence context

- ▶ Different types of co-occurrence (Evert 2008)
 - ▶ **surface context** (word or character window)
 - ▶ **textual context** (non-overlapping segments)
 - ▶ **syntactic context** (dependency relations)
 - 🔍 from research into collocations
- ▶ Context size
 - ▶ small context (few words, syntactic relation) → more specific
 - ▶ large context (many words, entire document) → more general
- ▶ Different roles of co-occurrence context
 - ▶ unstructured context → acts as a **filter** for counts
 - ▶ **structured** context → subcategorizes feature terms

🔍 What effects do you expect from these choices?

Surface context

Context term occurs **within a span of k words** around target.

The silhouette of the sun beyond a wide-open bay on the lake; the sun still glitters although evening has arrived in Kuhmo. It's midsummer; the living room has its instruments and other objects in each of its corners. [L3/R3 span, $k = 6$]

Parameters:

- ▶ span size (in words or characters)
- ▶ symmetric *vs.* one-sided span
- ▶ uniform or “triangular” (distance-based) weighting (don't!)
- ▶ spans clamped to sentences or other textual units?

Effect of span size

Nearest neighbours of *dog* (BNC)

2-word span

- ▶ cat
- ▶ horse
- ▶ fox
- ▶ pet
- ▶ rabbit
- ▶ pig
- ▶ animal
- ▶ mongrel
- ▶ sheep
- ▶ pigeon

30-word span

- ▶ kennel
- ▶ puppy
- ▶ pet
- ▶ bitch
- ▶ terrier
- ▶ rottweiler
- ▶ canine
- ▶ cat
- ▶ to bark
- ▶ Alsatian

<http://clit.cimec.unitn.it/infomap-query/>

Textual context

Context term is in the **same linguistic unit** as target.

The silhouette of the sun beyond a wide-open bay on the lake; the sun still glitters although evening has arrived in Kuhmo. It's midsummer; the living room has its instruments and other objects in each of its corners.

Parameters:

- ▶ choice of linguistic unit
 - ▶ sentence
 - ▶ paragraph
 - ▶ turn in a conversation
 - ▶ Web page
 - ▶ tweet

🔍 similar to large surface spans, but more self-contained

Syntactic context

Context term is linked to target by a **syntactic dependency** (e.g. subject, modifier, ...).

The **silhouette** of the **sun** beyond a wide-open **bay** on the lake; the **sun** still **glitters** although evening has arrived in Kuhmo. It's midsummer; the living room has its instruments and other objects in each of its corners.

Parameters:

- ▶ types of syntactic dependency (Padó & Lapata 2007)
- ▶ maximal length of dependency path (1 for direct relation)
- ▶ homogeneous data (e.g. only verb-object) *vs.* heterogeneous data (e.g. all children and parents of the verb)

Comparison of co-occurrence contexts

Contexts range from general/**implicit** to specific/**explicit**:

	features are
textual / large span	from same topic domain
small span	collocations
syntactic (single relation)	attributes (focus on aspect)
knowledge pattern	properties

“Knowledge pattern” context

Context term is linked to target by a **lexico-syntactic pattern** (text mining, cf. Hearst 1992, Pantel & Pennacchiotti 2008, etc.).

In Provence, Van Gogh painted with bright **colors** **such as** **red** **and** **yellow**. These **colors** **produce** incredible **effects** on anybody looking at his paintings.

Parameters:

- ▶ inventory of lexical patterns
 - ▶ lots of research to identify semantically interesting patterns (cf. Almuhareb & Poesio 2004, Veale & Hao 2008, etc.)
- ▶ fixed *vs.* flexible patterns
 - ▶ patterns are mined from large corpora and automatically generalised (optional elements, POS tags or semantic classes)

Structured vs. unstructured context

- ▶ In **unstructured** models, context specification acts as a **filter**
 - ▶ determines whether context token counts as co-occurrence
 - ▶ e.g. must be linked by any direct syntactic dependency relation
- ▶ In **structured** models, feature terms are **subtyped**
 - ▶ depending on their position in the context
 - ▶ e.g. left *vs.* right context, type of syntactic relation, etc.

Structured vs. unstructured surface context

A dog bites a man. The man's dog bites a dog. A dog bites a man.

unstructured	bite
dog	4
man	3

→ data are less sparse (L/R context aggregated)

A dog bites a man. The man's dog bites a dog. A dog bites a man.

structured	bite-L	bite-R
dog	1	3
man	2	1

→ more sensitive to semantic distinctions

Structured vs. unstructured dependency context

A dog bites a man. The man's dog bites a dog. A dog bites a man.

unstructured	bite
dog	4
man	2


→ data are less sparse (all syntactic relations aggregated)

A dog bites a man. The man's dog bites a dog. A dog bites a man.

structured	bite-subj	bite-obj
dog	3	1
man	0	2

→ more sensitive to semantic distinctions

Building a distributional model



Marginal and expected frequencies

► Matrix of **observed** co-occurrence frequencies not sufficient

target	feature	<i>O</i>	<i>R</i>	<i>C</i>	<i>E</i>
<i>dog</i>	<i>small</i>	855	33,338	490,580	134.34
<i>dog</i>	<i>domesticated</i>	29	33,338	918	0.25

► Notation

- *O* = observed co-occurrence frequency
- *R* = overall frequency of target term = row marginal frequency
- *C* = overall frequency of feature = column marginal frequency
- *N* = sample size ≈ size of corpus

► **Expected** co-occurrence **frequency** (cf. Evert 2008)

$$E = \frac{R \cdot C}{N} \longleftrightarrow O$$

Obtaining marginal frequencies (Evert 2008)

- ▶ Term-document matrix
 - ▶ R = frequency of target term in corpus
 - ▶ C = size of document (# tokens)
 - ▶ N = corpus size
- ▶ Syntactic co-occurrence
 - ▶ # of dependency instances in which target/feature participates
 - ▶ N = total number of dependency instances
 - ▶ N, R, C can be computed from full co-occurrence matrix M
- ▶ Textual co-occurrence
 - ▶ R, C, O are “document” frequencies, i.e. number of context units in which target, feature or combination occurs
 - ▶ N = total # of context units

Marginal frequencies in wordspace

DSM objects in wordspace (class `dsm`) include marginal frequencies as well as counts of nonzero cells for rows and columns.

```
> TT$rows
  term      f nnzero
1  cat  22007      5
2  dog  50807      7
3 animal 77053      7
4  time 115693      7
5 reason  95047      6
6  cause  54739      5
7 effect 133102      6
> TT$cols
...
> TT$globals$N
[1] 199902178
> TT$M # the full co-occurrence matrix
```

Obtaining marginal frequencies (Evert 2008)

- ▶ Surface co-occurrence
 - ▶ it is quite tricky to obtain fully consistent counts (Evert 2004)
 - ▶ **recommended**: correct E for span size k (= # tokens in span)¹


$$E = k \cdot \frac{R \cdot C}{N}$$

with R, C = individual corpus frequencies and N = corpus size

- ▶ can also be implemented by pre-multiplying $R' = k \cdot R$ (→ all pre-compiled surface DSMs in the course)
- ▶ alternatively, compute marginals and sample size by summing over full co-occurrence matrix (→ E as above, but inflated N)

¹NB: shifted PPMI (Levy & Goldberg 2014) corresponds to a post-hoc application of the span size adjustment. It performs worse than PPMI, but paper suggests they already approximate correct E by summing over matrix M .

Building a distributional model



Feature scaling

- ▶ **M** is often dominated by few very large entries
(→ highly skewed frequency distribution due to **Zipf's law**)
- ▶ Logarithmic scaling: $O' = \log(O + 1)$
(cf. Weber-Fechner law for human perception)
- ▶ Statistical **association measures** (Evert 2004, 2008) take frequency of target term and feature into account
 - ▶ usually based on comparison of observed and expected co-occurrence frequency
 - ▶ measures differ in how they balance O and E

Simple association measures

- ▶ pointwise **Mutual Information (MI)**

$$MI = \log_2 \frac{O}{E}$$

- ▶ local MI

$$\text{local-MI} = O \cdot MI = O \cdot \log_2 \frac{O}{E}$$

- ▶ **t-score**

$$t = \frac{O - E}{\sqrt{O}}$$

target	feature	O	E	MI	local-MI	t-score
dog	small	855	134.34	2.67	2282.88	24.64
dog	domesticated	29	0.25	6.85	198.76	5.34
dog	sgjkj	1	0.00027	11.85	11.85	1.00

Other association measures

- ▶ simple **log-likelihood** ($\approx$ local-MI)

$$G^2 = \pm 2 \cdot \left(O \cdot \log_2 \frac{O}{E} - (O - E) \right)$$

with positive sign for $O > E$ and negative sign for $O < E$

- ▶ **Dice** coefficient

$$\text{Dice} = \frac{2O}{R + C}$$

- ▶ Many other association measures (AMs) available, often based on full contingency tables (see Evert 2008)
 - ▶ <http://www.collocations.de/>
 - ▶ <http://sigil.r-forge.r-project.org/>

Applying association scores in wordspace

```
> options(digits=3) # print fractional values with limited precision
> dsm.score(TT, score="MI", sparse=FALSE, matrix=TRUE)
      breed  tail  feed  kill important explain likely
cat      6.21 4.568 3.129 2.801      -Inf 0.0182  -Inf
dog      7.78 3.081 3.922 2.323     -3.774 -1.1888 -0.4958
animal  3.50 2.132 4.747 2.832     -0.674 -0.4677 -0.0966
time    -1.65 -2.236 -0.729 -1.097    -1.728 -1.2382 0.6392
reason -2.30  -Inf -1.982 -0.388      1.472 4.0368 2.8860
cause   -Inf -0.834  -Inf -2.177      1.900 2.8329 4.0691
effect  -Inf -2.116 -2.468 -2.459      0.791 1.6312 0.9221
```

- ▶ sparseness of matrix representation is lost (try with TC!)
- ▶ cells with score $x = -\infty$ are inconvenient
- ▶ distribution of scores may be even more skewed than co-occurrence frequencies themselves (esp. for G^2)

Sparse association measures

- Sparse association scores are cut off at zero, i.e.

$$f(x) = \begin{cases} x & x > 0 \\ 0 & x \leq 0 \end{cases}$$

- Also known as “positive” scores
 - PPMI = positive pointwise MI (e.g. Bullinaria & Levy 2007)
 - wordspace computes sparse AMs by default → “MI” = PPMI
- Preserves sparseness if $x \leq 0$ for all empty cells ($O = 0$)
 - sparseness may even increase: cells with $x < 0$ become empty
- Further thinning may be beneficial (Polajnar & Clark 2014)
 - apply shifted cutoff threshold $x > \theta$ (Levy *et al.* 2015)
 - keep only k top-scoring features for each target

Score transformations

An additional scale transformation can be applied in order to de-skew association scores:


- signed **logarithmic** transformation

$$f(x) = \pm \log(|x| + 1)$$

- **sigmoid** transformation as soft binarization

$$f(x) = \tanh x$$

- **sparse** AM as (**shifted**) cutoff transformation (aka. ReLU)




Association scores & transformations in wordspace

```
> dsm.score(TT, score="MI", matrix=TRUE) # PPMI
      breed tail feed kill important explain likely
cat      6.21 4.57 3.13 2.80      0.000 0.0182 0.000
dog      7.78 3.08 3.92 2.32      0.000 0.0000 0.000
animal   3.50 2.13 4.75 2.83      0.000 0.0000 0.000
time      0.00 0.00 0.00 0.00      0.000 0.0000 0.639
reason    0.00 0.00 0.00 0.00      1.472 4.0368 2.886
cause     0.00 0.00 0.00 0.00      1.900 2.8329 4.069
effect    0.00 0.00 0.00 0.00      0.791 1.6312 0.922
> dsm.score(TT, score="simple-ll", matrix=TRUE)
> dsm.score(TT, score="simple-ll", transf="log", matrix=T)
# logarithmic co-occurrence frequency
> dsm.score(TT, score="freq", transform="log", matrix=T)

# now try other parameter combinations
> ?dsm.score # read help page for available parameter settings
```

Building a distributional model




Geometric distance = metric

- Distance between vectors $\mathbf{u}, \mathbf{v} \in \mathbb{R}^n \rightarrow$ (dis)similarity

- $\mathbf{u} = (u_1, \dots, u_n)$
- $\mathbf{v} = (v_1, \dots, v_n)$

- Euclidean** distance $d_2(\mathbf{u}, \mathbf{v})$
- "City block" **Manhattan** distance $d_1(\mathbf{u}, \mathbf{v})$
- Both are special cases of the **Minkowski** p -distance $d_p(\mathbf{u}, \mathbf{v})$ (for $p \in [1, \infty]$)



$$d_p(\mathbf{u}, \mathbf{v}) := (|u_1 - v_1|^p + \dots + |u_n - v_n|^p)^{1/p}$$


$$d_\infty(\mathbf{u}, \mathbf{v}) = \max\{|u_1 - v_1|, \dots, |u_n - v_n|\}$$

Geometric distance = metric

- Distance between vectors $\mathbf{u}, \mathbf{v} \in \mathbb{R}^n \rightarrow$ (dis)similarity

- $\mathbf{u} = (u_1, \dots, u_n)$
- $\mathbf{v} = (v_1, \dots, v_n)$

- Hamming** distance $d_0(\mathbf{u}, \mathbf{v})$ not very useful for DSM
- Extension of the Minkowski **p -distance** $d_p(\mathbf{u}, \mathbf{v})$ (for $0 \leq p \leq 1$)



$$d_p(\mathbf{u}, \mathbf{v}) := |u_1 - v_1|^p + \dots + |u_n - v_n|^p$$

$$d_0(\mathbf{u}, \mathbf{v}) = \#\{i \mid u_i \neq v_i\}$$

Computing distances

Preparation: store "scored" matrix in DSM object

```
> TT <- dsm.score(TT, score="freq", transform="log")
```

Compute distances between individual term pairs ...

```
> pair.distances(c("cat", "cause"), c("animal", "effect"),
  TT, method="euclidean")
cat/animal cause/effect
4.16 1.53
```

... or full distance matrix.


```
> dist.matrix(TT, method="euclidean")
> dist.matrix(TT, method="minkowski", p=4)
```

Distance and vector length = norm

- Intuitively, distance $d(\mathbf{u}, \mathbf{v})$ should correspond to length $\|\mathbf{u} - \mathbf{v}\|$ of displacement vector $\mathbf{u} - \mathbf{v}$

- $d(\mathbf{u}, \mathbf{v})$ is a **metric**
- $\|\mathbf{u} - \mathbf{v}\|$ is a **norm**
- $\|\mathbf{u}\| = d(\mathbf{u}, \mathbf{0})$

- Any norm-induced metric is **translation-invariant**
- Minkowski p -norm** with $d_p(\mathbf{u}, \mathbf{v}) = \|\mathbf{u} - \mathbf{v}\|_p$



$$\|\mathbf{u}\|_p := (|u_1|^p + \dots + |u_n|^p)^{1/p} \quad \text{for } 1 \leq p$$

$$\|\mathbf{u}\|_p := |u_1|^p + \dots + |u_n|^p \quad \text{for } 0 \leq p < 1 \text{ (an F-norm)}$$

$$\|\mathbf{u}\|_0 = \#\{i \mid u_i \neq 0\} \quad \|\mathbf{u}\|_\infty = \max\{|u_1|, \dots, |u_n|\}$$


Normalisation of row vectors

- Part 1: geometric distances only meaningful for vectors of the same length $\|\mathbf{x}\|$
- Normalize by scalar division:

$$\mathbf{x}' = \frac{1}{\|\mathbf{x}\|} \cdot \mathbf{x} = \left(\frac{x_1}{\|\mathbf{x}\|}, \frac{x_2}{\|\mathbf{x}\|}, \dots \right)$$

with $\|\mathbf{x}'\| = 1$

- Norm must be compatible with distance measure!
- Special case: scale $\mathbf{x} \geq 0$ to stochastic vector with $\|\mathbf{x}\|_1 = |x_1| + \dots + |x_n|$
→ probabilistic interpretation



Norms and normalization

```
> rowNorms(TT$S, method="euclidean")
  cat  dog animal  time reason  cause effect
6.90  8.96  8.82 10.29  8.13  6.86  6.52
```

```
> TT <- dsm.score(TT, score="freq", transform="log",
                  normalize=TRUE, method="euclidean")
> rowNorms(TT$S, method="euclidean") # all = 1 now
> dist.matrix(TT, method="euclidean")
  cat  dog animal  time reason  cause effect
cat  0.000 0.224  0.473 0.782  1.121 1.239  1.161
dog  0.224 0.000  0.398 0.698  1.065 1.179  1.113
animal 0.473 0.398  0.000 0.426  0.841 0.971  0.860
time  0.782 0.698  0.426 0.000  0.475 0.585  0.502
reason 1.121 1.065  0.841 0.475  0.000 0.277  0.198
cause  1.239 1.179  0.971 0.585  0.277 0.000  0.224
effect 1.161 1.113  0.860 0.502  0.198 0.224  0.000
```

Distance measures for non-negative vectors

- Information theory: **Kullback-Leibler (KL) divergence** for stochastic vectors (non-negative $\mathbf{x} \geq 0$ and $\|\mathbf{x}\|_1 = 1$)

$$D(\mathbf{u} \parallel \mathbf{v}) = \sum_{i=1}^n u_i \cdot \log_2 \frac{u_i}{v_i}$$

- Properties of KL divergence
 - most appropriate for a probabilistic interpretation of $\mathbf{M}$
 - zeroes in $\mathbf{v}$ without corresponding zeroes in $\mathbf{u}$ are problematic
 - not symmetric**, unlike geometric distance measures
 - alternatives: skew divergence, Jensen-Shannon divergence
- A symmetric distance metric (Endres & Schindelin 2003)


$$D_{uv} = D(\mathbf{u} \parallel \mathbf{z}) + D(\mathbf{v} \parallel \mathbf{z}) \quad \text{with} \quad \mathbf{z} = \frac{\mathbf{u} + \mathbf{v}}{2}$$

Similarity measures

- Angle α between vectors $\mathbf{u}, \mathbf{v} \in \mathbb{R}^n$ is given by

$$\begin{aligned} \cos \alpha &= \frac{\sum_{i=1}^n u_i \cdot v_i}{\sqrt{\sum_i u_i^2} \cdot \sqrt{\sum_i v_i^2}} \\ &= \frac{\mathbf{u}^T \mathbf{v}}{\|\mathbf{u}\|_2 \cdot \|\mathbf{v}\|_2} \end{aligned}$$

- cosine** measure of similarity: $\cos \alpha$
 - $\cos \alpha = 1 \rightarrow$ collinear
 - $\cos \alpha = 0 \rightarrow$ orthogonal
- Corresponding metric: **angular distance** α



Euclidean distance or cosine similarity?

$$\begin{aligned}
 d_2(\mathbf{u}, \mathbf{v}) &= \|\mathbf{u} - \mathbf{v}\|_2 = \sqrt{\sum_i (u_i - v_i)^2} \\
 &= \sqrt{\sum_i u_i^2 + \sum_i v_i^2 - 2 \sum_i u_i v_i} \\
 &= \sqrt{\|\mathbf{u}\|_2^2 + \|\mathbf{v}\|_2^2 - 2 \mathbf{u}^T \mathbf{v}} \\
 &= \sqrt{2 - 2 \cos \phi}
 \end{aligned}$$

☞ $d_2(\mathbf{u}, \mathbf{v})$ is a monotonically increasing function of ϕ

Euclidean distance and cosine similarity are equivalent: if vectors have been normalised ($\|\mathbf{u}\|_2 = \|\mathbf{v}\|_2 = 1$), both lead to the same neighbour ranking.

Similarity measures for non-negative vectors

- Generalized **Jaccard coefficient** = shared features


$$J(\mathbf{u}, \mathbf{v}) = \frac{\sum_{i=1}^n \min\{u_i, v_i\}}{\sum_{i=1}^n \max\{u_i, v_i\}}$$

- $1 - J(\mathbf{u}, \mathbf{v})$ is a distance **metric** (Kosub 2016)

- An asymmetric measure of feature **overlap** (Clarke 2009)

$$o(\mathbf{u}, \mathbf{v}) = \frac{\sum_{i=1}^n \min\{u_i, v_i\}}{\sum_{i=1}^n u_i}$$

Building a distributional model



Dimensionality reduction = model compression

- Co-occurrence matrix **M** is often unmanageably large and can be extremely sparse
 - Google Web1T5: $1\text{M} \times 1\text{M}$ matrix with one trillion cells, of which less than 0.05% contain nonzero counts (Evert 2010)
- Compress matrix by reducing dimensionality (= rows)
- **Feature selection**: columns with high frequency & variance
 - measured by entropy, chi-squared test, nonzero count, ...
 - may select similar dimensions and discard valuable information
- **Projection** into (linear) subspace
 - principal component analysis (PCA)
 - independent component analysis (ICA)
 - random indexing (RI)
 - ☞ intuition: preserve distances between data points


Dimensionality reduction & latent dimensions

Landauer & Dumais (1997) claim that LSA dimensionality reduction (and related PCA technique) uncovers **latent dimensions** by exploiting correlations between features.

- ▶ Example: term-term matrix
- ▶ V-Obj co-oc. extracted from BNC
 - ▶ targets = noun lemmas
 - ▶ features = verb lemmas
- ▶ feature scaling: association scores (SketchEngine log Dice)
- ▶ $k = 186$ nouns with $f_{\text{buy}} + f_{\text{sell}} \geq 25$
- ▶ $n = 2$ dimensions: *buy* and *sell*

noun	buy	sell
<i>antique</i>	5.12	5.50
<i>bread</i>	5.96	3.99
<i>computer</i>	6.75	6.83
<i>factory</i>	4.95	4.72
<i>group</i>	4.93	4.28
<i>jewellery</i>	5.11	5.73
<i>mill</i>	5.14	5.41
<i>people</i>	3.00	4.26
<i>record</i>	6.81	6.68
<i>souvenir</i>	5.45	4.67
<i>ticket</i>	8.93	8.74


Dimensionality reduction & latent dimensions




Motivating latent dimensions & subspace projection

- ▶ The **latent property** of being a commodity is “expressed” through associations with several verbs: *sell*, *buy*, *acquire*, ...
- ▶ Consequence: these DSM dimensions will be **correlated**
- ▶ Identify **latent dimension** by looking for strong correlations (or weaker correlations between large sets of features)
- ▶ Projection into subspace V of $k < n$ latent dimensions as a “**noise reduction**” technique → **LSA**
- ▶ Assumptions of this approach:
 - ▶ “latent” distances in V are semantically meaningful
 - ▶ other “residual” dimensions represent chance co-occurrence patterns, often particular to the corpus underlying the DSM


PCA dimensionality reduction: orthogonal projection




PCA dimensionality reduction: orthogonal projection



PCA dimensionality reduction: orthogonal projection



PCA dimensionality reduction: further dimensions



PCA dimensionality reduction

- ▶ Principal component analysis (**PCA**)
 - ▶ orthogonal projection into orthogonal latent dimensions
 - ▶ finds optimal subspace of given dimensionality (such that orthogonal projection preserves distance information)
 - ▶ but requires features centered at 0 → no longer sparse
- ▶ Singular value decomposition (**SVD**)
 - ▶ the mathematical algorithm behind PCA
 - ▶ often applied without centering in distributional semantics
 - ⚠ optimality of subspace not guaranteed
 - ⚠ first dimension(s) uninteresting (↔ non-negative quadrant)
- ▶ NB: row vectors should be renormalised after PCA/SVD
 - ▶ unless cosine similarity / angular distance is used
 - ⚠ also normalise vectors **before** dimensionality reduction


Dimensionality reduction in practice

```
# SVD is the algorithm behind PCA dimensionality reduction
> TT2 <- dsm.projection(TT, n=2, method="svd")
> TT2
      svd1    svd2
cat   -0.733 -0.6615
dog   -0.782 -0.6110
animal -0.914 -0.3606
time  -0.993  0.0302
reason -0.889  0.4339
cause  -0.817  0.5615
effect -0.871  0.4794

> x <- TT2[, 1] # first latent dimension
> y <- TT2[, 2] # second latent dimension
> plot(x, y, pch=20, col="red",
      xlim=extendrange(x), ylim=extendrange(y))
> text(x, y, rownames(TT2), pos=3)
```

Scaling latent dimensions

- ▶ Truncated SVD omits latent dimensions that capture relatively little distance information (here $r = 400$)
- ▶ Skip first k dimensions, e.g. $k = 50$ (Bullinaria & Levy 2012)
- ▶ Power-scaling of dimensions: σ^P (Caron 2001)
 - ▶ Bullinaria & Levy (2012) report positive effect
 - ▶ esp. with $P = 0$ to equalize dimensions (**whitening**)




Power-scaling in practice

```
> TT2 <- dsm.projection(TT, n=2, method="svd", power=0)
> TT2
      svd1    svd2
cat   -0.322 -0.5110
dog   -0.343 -0.4721
animal -0.401 -0.2786
time  -0.436  0.0233
reason -0.390  0.3353
cause  -0.359  0.4338
effect -0.383  0.3704

# power-scaling can also be applied post-hoc
> sigma <- attr(TT2, "sigma") # singular values
> scaleMargins(TT2, cols=sigma^0.5) # P = 1/2
> scaleMargins(TT2, cols=sigma) # unscaled (P = 1)
```

Dimensionality reduction by RI

- ▶ Random indexing (**RI**)
 - ▶ project into random subspace (Sahlgrén & Karlgrén 2005)
 - ▶ reasonably good if there are many subspace dimensions
 - ▶ can be performed online w/o collecting full co-oc. matrix



Outline

DSM parameters

A taxonomy of DSM parameters
 Context type & size
 Feature scaling
 Measuring distance
 Dimensionality reduction
 Examples

Building a DSM

Sparse matrices
 Example: a verb-object DSM

Some well-known DSM examples

Latent Semantic Analysis (Landauer & Dumais 1997)

- ▶ term-context matrix with document context
- ▶ weighting: log term frequency and term entropy
- ▶ distance measure: cosine
- ▶ dimensionality reduction: SVD

Hyperspace Analogue to Language (Lund & Burgess 1996)

- ▶ term-term matrix with surface context
- ▶ structured (left/right) and distance-weighted frequency counts
- ▶ distance measure: Minkowski metric ($1 \leq p \leq 2$)
- ▶ dimensionality reduction: feature selection (high variance)

Some well-known DSM examples

Infomap NLP (Widdows 2004)

- ▶ term-term matrix with unstructured surface context
- ▶ weighting: none
- ▶ distance measure: cosine
- ▶ dimensionality reduction: SVD

Random Indexing (Karlgrén & Sahlgrén 2001)

- ▶ term-term matrix with unstructured surface context
- ▶ weighting: various methods
- ▶ distance measure: various methods
- ▶ dimensionality reduction: random indexing (RI)

Some well-known DSM examples

Dependency Vectors (Padó & Lapata 2007)

- ▶ term-term matrix with unstructured dependency context
- ▶ weighting: log-likelihood ratio
- ▶ distance measure: PPMI-weighted Dice (Lin 1998)
- ▶ dimensionality reduction: none

Distributional Memory (Baroni & Lenci 2010)

- ▶ term-term matrix with structured and unstructured dependencies + knowledge patterns
- ▶ weighting: local-MI on type frequencies of link patterns
- ▶ distance measure: cosine
- ▶ dimensionality reduction: none

... and an unexpected application

Authorship attribution (Burrows 2002)

- ▶ Burrows's Delta method is very popular in modern literary stylometry and authorship attribution (Evert *et al.* 2017)
- ▶ document-term matrix with word forms as features
- ▶ weighting: relative frequency of word form in document
- ▶ feature selection: 200–5,000 most frequent words (mfw)
- ▶ columns are standardized ($\mu = 0$, $\sigma^2 = 1$) → z-scores
- ▶ clustering of documents based on various distance metrics (or nearest-neighbour classifier for known authors)
- ▶ dimensionality reduction: none
- ▶ main result: angle/cosine $\succ$ Manhattan $\succ$ Euclidean

Outline

DSM parameters

A taxonomy of DSM parameters

Context type & size

Feature scaling

Measuring distance

Dimensionality reduction

Examples

Building a DSM

Sparse matrices

Example: a verb-object DSM

Scaling up to the real world

- ▶ So far, we have worked on minuscule **toy models**
- 🗨️ We want to scale up to **real world** data sets now
- ▶ Example 1: span-based DSM on BNC content words
 - ▶ 83,926 lemma types with $f \geq 10$
 - ▶ term-term matrix with $83,926 \cdot 83,926 = 7$ billion entries
 - ▶ standard representation requires 56 GB of RAM (8-byte floats)
 - ▶ only 22.1 million non-zero entries (= 0.32%)
- ▶ Example 2: Google Web 1T 5-grams (1 trillion words)
 - ▶ more than 1 million word types with $f \geq 2500$
 - ▶ term-term matrix with 1 trillion entries requires 8 TB RAM
 - ▶ only 400 million non-zero entries (= 0.04%)

Sparse matrix representation

- ▶ Invented example of a **sparingly populated** DSM matrix

	eat	get	hear	kill	see	use
boat	.	59	.	.	39	23
cat	.	.	.	26	58	.
cup	.	98	.	.	.	.
dog	33	.	42	.	83	.
knife	.	.	.	.	.	84
pig	9	.	.	27	.	.

- ▶ Store only non-zero entries in compact **sparse matrix format**

row	col	value	row	col	value
1	2	59	4	1	33
1	5	39	4	3	42
1	6	23	4	5	83
2	4	26	5	6	84
2	5	58	6	1	9
3	2	98	6	4	27

Working with sparse matrices

- ▶ Compressed format: each row index (or column index) stored only once, followed by non-zero entries in this row (or column)
 - ▶ convention: **column-major** matrix (data stored by columns)
- ▶ Specialised algorithms for sparse matrix algebra
 - ▶ especially matrix multiplication, solving linear systems, etc.
 - ▶ take care to avoid operations that create a dense matrix!
- ▶ R implementation: **Matrix** package
 - ▶ essential for real-life distributional semantics
 - ▶ **wordspace** provides additional support for sparse matrices (vector distances, sparse SVD, ...)
- ▶ Other software: Matlab, Octave, Python + SciPy
 - ▶ TensorFlow, PyTorch, ... always use dense matrices!

Outline

DSM parameters

A taxonomy of DSM parameters
 Context type & size
 Feature scaling
 Measuring distance
 Dimensionality reduction
 Examples

Building a DSM

Sparse matrices

Example: a verb-object DSM

Triplet tables

- ▶ A sparse DSM matrix can be represented as a table of triplets (target, feature, co-occurrence frequency)
 - ▶ for syntactic co-occurrence and term-document matrices, marginals can be computed from a complete triplet table
 - ▶ for surface and textual co-occurrence, marginals have to be provided in separate files (see `?read.dsm.triplet`)

noun	rel	verb	f	mode
dog	subj	bite	3	spoken
dog	subj	bite	12	written
dog	obj	bite	4	written
dog	obj	stroke	3	written
...	...	...	...	...

- ▶ `DSM_VerbNounTriples_BNC` contains additional information
 - ▶ syntactic relation between noun and verb
 - ▶ written or spoken part of the British National Corpus

Constructing a DSM from a triplet table

- ▶ Additional information can be used for filtering (verb-object relation), or aggregate frequencies (spoken + written BNC)

```
> tri <- subset(DSM_VerbNounTriples_BNC, rel == "obj")
```

- ▶ Construct DSM object from triplet input
 - ▶ `raw.freq=TRUE` indicates raw co-occurrence frequencies (rather than a pre-weighted DSM)
 - ▶ constructor aggregates counts from duplicate entries
 - ▶ marginal frequencies are automatically computed

```
> VObj <- dsm(target=tri$noun, feature=tri$verb,
              score=tri$f, raw.freq=TRUE)
> VObj # inspect marginal frequencies (e.g. head(VObj$rows, 20))
```

Exploring the DSM

```
> VObj <- dsm.score(VObj, score="MI", normalize=TRUE)

> nearest.neighbours(VObj, "dog") # angular distance
  horse    cat  animal  rabbit    fish    guy
  73.9    75.9    76.2    77.0    77.2    78.5
  cichlid  kid    bee creature
  78.6    79.0    79.1    79.5

> nearest.neighbours(VObj, "dog", method="manhattan")
# NB: we used an incompatible Euclidean normalization!

> VObj50 <- dsm.projection(VObj, n=50, method="svd")
> nearest.neighbours(VObj50, "dog")
```

References I

- Baroni, Marco and Lenci, Alessandro (2010). Distributional Memory: A general framework for corpus-based semantics. *Computational Linguistics*, **36**(4), 673–712.
- Bullinaria, John A. and Levy, Joseph P. (2007). Extracting semantic representations from word co-occurrence statistics: A computational study. *Behavior Research Methods*, **39**(3), 510–526.
- Bullinaria, John A. and Levy, Joseph P. (2012). Extracting semantic representations from word co-occurrence statistics: Stop-lists, stemming and SVD. *Behavior Research Methods*, **44**(3), 890–907.
- Burrows, John (2002). 'Delta': a measure of stylistic difference and a guide to likely authorship. *Literary and Linguistic Computing*, **17**(3), 267–287.
- Caron, John (2001). Experiments with LSA scoring: Optimal rank and basis. In M. W. Berry (ed.), *Computational Information Retrieval*, pages 157–169. Society for Industrial and Applied Mathematics, Philadelphia, PA, USA.
- Clarke, Daoud (2009). Context-theoretic semantics for natural language: an overview. In *Proceedings of the Workshop on Geometrical Models of Natural Language Semantics*, pages 112–119, Athens, Greece.
- Endres, Dominik M. and Schindelin, Johannes E. (2003). A new metric for probability distributions. *IEEE Transactions on Information Theory*, **49**(7), 1858–1860.

Practice

- ▶ Code examples and further explanations: [hands_on_part2.R](#)
- ▶ How many different models can you build from [DSM_VerbNounTriples_BNC?](#)
 - ▶ apply different filters, scores, transformations and metrics
 - ▶ explore nearest neighbours of selected word
- ▶ Build real-life DSMs from pre-compiled co-occurrence data
 - ▶ <http://wordspace.collocations.de/doku.php/course:material>
 - ▶ load pre-compiled matrix and apply different parameters
 - ▶ compare nearest neighbours or semantic maps
- ▶ Learn how to import your own co-occurrence data
 - ▶ [hands_on_part2_input_formats.R](#)
 - ▶ download example data sets to subdirectory [data/](#)

References II

- Evert, Stefan (2004). *The Statistics of Word Cooccurrences: Word Pairs and Collocations*. Dissertation, Institut für maschinelle Sprachverarbeitung, University of Stuttgart.
- Evert, Stefan (2008). Corpora and collocations. In A. Lüdeling and M. Kytö (eds.), *Corpus Linguistics. An International Handbook*, chapter 58, pages 1212–1248. Mouton de Gruyter, Berlin, New York.
- Evert, Stefan (2010). Google Web 1T5 n-grams made easy (but not for the computer). In *Proceedings of the 6th Web as Corpus Workshop (WAC-6)*, pages 32–40, Los Angeles, CA.
- Evert, Stefan; Proisl, Thomas; Jannidis, Fotis; Reger, Isabella; Pielström, Steffen; Schöch, Christof; Vitt, Thorsten (2017). Understanding and explaining Delta measures for authorship attribution. *Digital Scholarship in the Humanities*, **22**(suppl_2), ii4–ii16.
- Karlgren, Jussi and Sahlgren, Magnus (2001). From words to understanding. In Y. Uesaka, P. Kanerva, and H. Asoh (eds.), *Foundations of Real-World Intelligence*, chapter 294–308. CSLI Publications, Stanford.
- Kosub, Sven (2016). A note on the triangle inequality for the Jaccard distance. [arXiv:1612.02696](#).

References III

- Landauer, Thomas K. and Dumais, Susan T. (1997). A solution to Plato's problem: The latent semantic analysis theory of acquisition, induction and representation of knowledge. *Psychological Review*, **104**(2), 211–240.
- Levy, Omer and Goldberg, Yoav (2014). Neural word embedding as implicit matrix factorization. In *Proceedings of Advances in Neural Information Processing Systems 27*, pages 2177–2185. Curran Associates, Inc.
- Levy, Omer; Goldberg, Yoav; Dagan, Ido (2015). Improving distributional similarity with lessons learned from word embeddings. *Transactions of the Association for Computational Linguistics*, **3**, 211–225.
- Lin, Dekang (1998). Automatic retrieval and clustering of similar words. In *Proceedings of the 17th International Conference on Computational Linguistics (COLING-ACL 1998)*, pages 768–774, Montreal, Canada.
- Lund, Kevin and Burgess, Curt (1996). Producing high-dimensional semantic spaces from lexical co-occurrence. *Behavior Research Methods, Instruments, & Computers*, **28**(2), 203–208.
- Padó, Sebastian and Lapata, Mirella (2007). Dependency-based construction of semantic space models. *Computational Linguistics*, **33**(2), 161–199.

References IV

- Polajnar, Tamara and Clark, Stephen (2014). Improving distributional semantic vectors through context selection and normalisation. In *Proceedings of the 14th Conference of the European Chapter of the Association for Computational Linguistics*, pages 230–238, Gothenburg, Sweden.
- Sahlgren, Magnus and Karlgren, Jussi (2005). Automatic bilingual lexicon acquisition using random indexing of parallel corpora. *Natural Language Engineering*, **11**, 327–341.
- Widdows, Dominic (2004). *Geometry and Meaning*. Number 172 in CSLI Lecture Notes. CSLI Publications, Stanford.